

Social Media as a Measurement Tool of Depression in Populations

Munmun De Choudhury Scott Counts Eric Horvitz

Microsoft Research

One Microsoft Way, Redmond WA 98052

munmund@microsoft.com counts@microsoft.com horvitz@microsoft.com

ABSTRACT

Depression is a serious and widespread public health challenge. We examine the potential for leveraging social media postings as a new type of lens in understanding depression in populations. Information gleaned from social media bears potential to complement traditional survey techniques in its ability to provide finer grained measurements over time while radically expanding population sample sizes. We present work on using a crowdsourcing methodology to build a large corpus of postings on Twitter that have been shared by individuals diagnosed with clinical depression. Next, we develop a probabilistic model trained on this corpus to determine if posts could indicate depression. The model leverages signals of social activity, emotion, and language manifested on Twitter. Using the model, we introduce a social media depression index that may serve to characterize levels of depression in populations. Geographical, demographic and seasonal patterns of depression given by the measure confirm psychiatric findings and correlate highly with depression statistics reported by the Centers for Disease Control and Prevention (CDC).

Author Keywords

behavior, depression, emotion, health, language, social media, mental health, public health, Twitter, wellness

ACM Classification Keywords

H.5.m [Information Systems]: Information Systems Applications – Miscellaneous

INTRODUCTION

Depression affects more than 27 million Americans and is believed to be responsible for the more than 30,000 suicides each year [2,14]. Besides being directly debilitating to sufferers, depression can adversely affect chronic health conditions, such as cardiovascular disease, cancer, diabetes, and obesity. It is also known to have negative influences on individuals' family and personal relationships, work or school life, and sleeping and eating habits.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

WebSci'13, May 2–4, 2013, Paris, France.

Copyright 2013 ACM 978-1-4503-1889-1....\$10.00.

Over the coming 20 years, depression is projected to be the leading cause of disability in high-income nations such as the United States [16]. The World Health Organization (WHO) now ranks major depression as one of the most burdensome diseases in the world [2,16]. Although a number of primary care programs have been devised for its detection and treatment, the majority of the millions of Americans who meet depression criteria are untreated or undertreated [11]. Furthermore, ethnic minority groups such as Mexican Americans and African Americans are significantly less likely to receive depression therapies than are other ethnic groups [9].

As part of a national-scale effort to curb depression, every few years the Centers for Disease Control and Prevention (CDC) administers the Behavioral Risk Factor Surveillance System (BRFSS) survey via telephone to estimate the rate of depression among adults in the US [2]. However the large temporal gaps across which these measurements are made, as well as the limited number of participant responses (on the order of thousands) makes it difficult for agencies to track and identify risk factors that may be associated with mental illness, or to develop effective intervention programs.

We examine the potential of social media as a new tool for mental health measurement and surveillance. Platforms such as Twitter and Facebook are increasingly gaining traction among individuals allowing them to share their thoughts and emotions around a variety of happenings in everyday life. The emotion and language used in social media postings may indicate feelings of worthlessness, guilt, helplessness, and self-hatred that characterize depression as manifested in everyday life. Additionally, depression sufferers often show withdrawal from social situations and activities—i.e., the etiology of depression typically includes social environmental factors [17]. Characterization of social media activity and changing social ties within social media can provide measurement of such withdrawal and capture the depression sufferers' social context in a manner that might help detect depression in populations.

Relying on social media as a behavioral health assessment tool has other advantages as well. For instance, in contrast to the self-report methodology in behavioral surveys, where responses are prompted by the experimenter and typically comprise recollection of (sometimes subjective) health

facts, social media measurement of behavior captures social activity and language expression in a naturalistic setting. Such activity is real-time, and happens in the course of a person's day-to-day life. Hence it is less vulnerable to memory bias or experimenter demand effects, and can help track concerns at a fine-grained temporal scale.

Our main contributions in this paper are as follows:

(1) Using crowdsourcing techniques, we gather a ground truth set of 69K Twitter postings shared by individuals suffering from clinical depression—depression was measured using the CES-D (Center for Epidemiologic Studies Depression Scale) screening test [22].

(2) We develop statistical models (an SVM classifier) that can predict whether or not a Twitter post in a test set could be depression-indicative. To construct and test the predictive models, we harness evidence from a variety of measures, spanning emotional expression, linguistic style, user engagement, and egocentric social network properties. We demonstrate that our models can predict if a post is depression-indicative, with accuracy of more than 70% and precision of 0.82.

(3) Finally we propose a metric we refer to as the *social media depression index* (SMDI). SMDI uses the above prediction models to determine depressive-indicative postings on Twitter, and thereby helps characterize the levels of depression in populations. We conduct a variety of analyses at population scale, examining depression levels (as given by SMDI) across geography (US cities and states), demographics (gender), and time, including diurnal and seasonal patterns. Our findings from these analyses align with CDC reported statistics of depression in US population, as well as confirm known characteristics of depression given in clinical literature.

We believe that, when tied to behavioral health records from agencies, information derived from our prediction models and analyses can be valuable to epidemiologists who study macro-trends of individuals suffering from depression or other types of mental health disorders.

RELATED WORK

Social/Psychological Context and Depression

Offline social networks and attributes relating to the psychological environment of individuals have consistently been used to study behavioral health concerns. Billings and Moos [1] studied the roles of stress, social resources, and coping among individuals entering treatment for depression. Neils Rosenquist, Fowler, and Christakis [17] found that levels of depression showed diffusion upto three degrees of separation in a large social network, suggesting a network influence component to depression. On similar lines, in [10], Kawachi et al. explored the role of social ties and social capital in the maintenance of psychological wellbeing and treatment of behavioral health concerns. This prior research provides strong evidence that individuals'

social environments contain vital information useful for understanding and intervening on mental health.

In the field of psycholinguistics, Oxman et al. [18] demonstrated that linguistic analysis of speech could classify patients into diagnostic groups such as those suffering depression and paranoia. Computerized analysis of written text has also been known to reveal cues about neurotic tendencies and psychiatric disorders [23]. Utilizing such analyses, particularly of social media given their strong connection to people's social environment, can help us overcome the limitations of surveys for understanding the social/psychological context of individuals.

Public Health using Online Data

Leveraging internet data for modeling and analyzing public health behaviors has been a ripe area of research in the recent past. Google Flu Trends¹ provides nuanced predictions of flu infections based on online search queries. Paul and Dredze [20] developed a disease-specific topic model based on Twitter's posts in order to model behavior around a variety of diseases of importance in public health. Through language modeling of Twitter posts, Collier et al. [4] found evidence of high correlation between social media signals and diagnostic influenza case data. Sadelik et al. [24] developed statistical models that predicted infectious disease (e.g. flu) spread in individuals based on geotagged postings made on Twitter (also see [13]).

While this body of work has investigated a range of challenges around public health, research on harnessing social media for understanding behavioral health disorders is still in its infancy. Park et al. [19] found initial evidence that people do post about their depression and even their treatment for depression on Twitter. In other related work [6], we examined linguistic and emotional correlates for postnatal course of new mothers, and thereafter built a model to predict extreme behavioral changes in new mothers. This early work thus points to the potential of social media as a signal to leverage in the study of depression. With the present work we expand the scope of social media-based studies of depression by examining general depression at population scale.

DATA

A primary challenge is the gathering of adequate social media data, in particular Twitter posts, which bear information on whether or not they could contain depression-bearing/indicative content. We tackle this issue by collecting actual depression information on a set of Twitter users. The idea is that the postings made by a user suffering from depression will contain cues or signals that could reflect their mental condition. Such postings would thus help construct the positive/target class (i.e. set of posts that are depression-indicative) necessary for classification.

¹ <http://www.google.org/flutrends/>

Identifying Clinically Depressed Users

We employ a crowdsourcing methodology to collect information on the status of depression of a set of users.

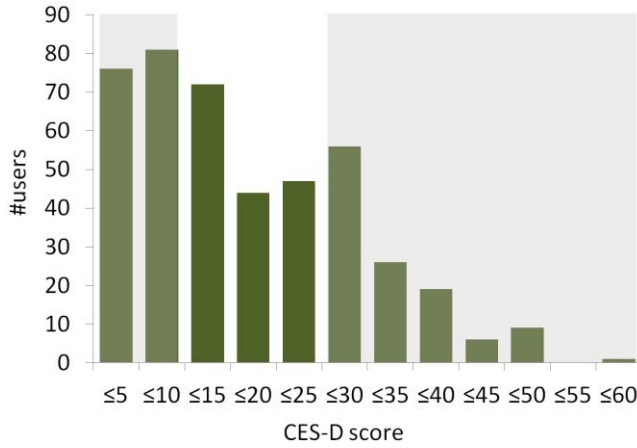


Figure 1. Distribution of depression scores based on the CES-D questionnaire. Shaded regions show the users who were considered for constructing ground truth classes.

Using Amazon’s Mechanical Turk interface, we designed HITs wherein crowdworkers were asked to take a clinical depression survey, followed by self-report seeking questions on their depression history. The crowdworkers could also opt-in to share their Twitter usernames if they had a *public* profile. We sought responses from crowdworkers who were located in the US, and had an approval rating greater than or equal to 90%. Each crowdworker was restricted to take the HIT exactly once, and was paid 90 cents for completing the task.

We used the CES-D (Center for Epidemiologic Studies Depression Scale) questionnaire as the primary tool to determine the depression levels in the crowdworkers. The CES-D is a 20-item self-report scale designed to measure depressive symptoms in the general population [22]. Higher scores in the test indicate the presence of more symptomatology. In our study, we also collected information about the crowdworkers’ depression history through the following questions:

- *Whether or not they have been diagnosed with clinical depression in the past. If so, when.*
- *If the answer above was positive, what was the estimated date of onset of depression.*
- *From the date of depression onset, how many depression episodes had they experienced.*

Finally, we asked the crowdworkers to share their (public) Twitter username for research purposes, under the privacy clause that their data may be mined using a computer program and analyzed anonymously.

Statistics of Crowdsourcing Study

A total of 1,583 crowdworkers completed our HITs between September 15 and October 31, 2012. However, not all crowdworkers opted in to share their Twitter

usernames—637 participants (~40%) agreed to provide us access to their Twitter feeds. Next, we eliminated noisy respondents who took too little time to complete the task.

Finally, we sought data from individuals with depression onset dates anytime in the last one year, but no later than three months before the survey was taken. These constraints were assumed to ensure that, for each user detected with depression, we would have sufficient postings on Twitter after the onset and until the point in time that they took the survey. We also sought users who had reported to be clinically depressed (and were currently so as well based on the scoring on the survey) with at least two depression episodes in the studied period. This helped us to focus on high precision data for which we had recurring signs of depression; depression would be a *prevalent* condition in these individuals. A set of 489 users was obtained, who indicated either they have clinical depression with onset in or after September 2011 and before June 2012, or did not have depression anytime. The set contained 251 males, 238 females, and the median age was 25 years.

Total number of users	489
(1) #users detected with depression	117
(2) #users with no signs of depression	157
Total number of Twitter posts	69,514
(1) Positive class (depression-indicative)	23,984
(2) Negative class (standard posts)	45,530
Mean number of posts per user	253.7
Variance of number of posts per user	24.88
Mean number of posts per day per user	2.79
Variance of number of posts per day per user	7.04

Table 1. Statistics of Twitter data for classification.

A distribution of the users per their depression scores is displayed in Figure 1. Based on the standard cut-off of 30 or above for high-range depression [21], we found 117 users with signs of severe depression. Separately, we obtained a set of 157 users, with scores in the range (0-10), who showed very little likelihood/almost no sign of depression in the period of our interest.

Positive class (depression-indicative posts)
“Are you okay?” Yes.... I understand that I am upset and hopeless and nothing can help me... I’m okay... but I am not alright
“empty” feelings I WAS JUST TALKING ABOUT HOW I I HAVE EMOTION OH MY GOODNESS I FEEL AWFUL
I want someone to hold me and be there for me when I’m sad.
Reloading twitter till I pass out. *lonely* *anxious* *butthurt* *frustrated* *dead*
Having a job again makes me happy. Less time to be depressed and eat all day while watching sad movies.

Table 2. Example posts from the positive class.

Building Ground Truth Dataset

We now discuss how we constructed a dataset of Twitter posts, with ground-truth label information (on whether or not the post is depression-indicative). We use Twitter

postings of the two sets of users obtained above—(1) for the positive class (*depression-indicative* posts), and (2) for the negative class (*standard* posts). For each user in both sets, we use all of their postings in a three-month period. For the positive class of posts, we collect all postings made by users in the first set during the three months after their indicated onset of depression. For the negative class, we collect all postings made by users in the second set, during the three months before the date they took our survey. For collecting the posts, we utilized the Firehose made available to us via our organization’s contract with Twitter. Table 1 gives the final statistics of the dataset thus constructed. A few examples of posts from the target class of depression-indicative posts are given in Table 2.

FEATURES

We propose several features to characterize the postings in our dataset. The features can be categorized into two types: post-centric and user-centric—the former captures properties in the post, while the latter characterizes the behavior of the post’s author. Several of our features are motivated from [6], where greater details can be accessed by the readers.

Post Features

Emotion. We consider four features of the emotional state manifested in the posts: *positive affect* (PA), *negative affect* (NA), *activation*, and *dominance*. Measurements of PA and NA per post are computed using the psycholinguistic resource LIWC (<http://www.liwc.net/>), whose emotion categories have been validated to perform well for determining affect in Twitter [3,6]. We use the ANEW lexicon [5,6] for computing activation and dominance.

Time. We define a measure that uses the timestamp information of a post, that is, whether it was a day time or night time post (in terms of local time of the post author). Our motivation springs from observations in the depression literature indicating that users showing depression signs tend to be relatively more active during the evening and night [15]. For the purpose we define a “night” window in a given day as “9PM—5:59AM” (consequently the “day” window for the same user, in local time, would be “6AM-8:59PM”). For each post, we therefore assign a “time” index 1 or -1, depending on whether it was posted respectively in the night or day window.

Linguistic Style. We also introduce features to characterize posts based on the use of linguistic styles [3,6]. We again use LIWC for determining 22 specific linguistic styles: *articles*, *auxiliary verbs*, *conjunctions*, *adverbs*, *impersonal pronouns*, *personal pronouns*, *prepositions*, *functional words*, *fillers*, *assent*, *negation*, *certainty* and *quantifiers*.

***n*-grams.** We also extract unigrams and bigrams from the posts, in order to account for the general language use.

User Features

Engagement. We utilize a set of engagement measures of the authors of posts in order to characterize the general

behavior associated with them—we believe users found to be clinically depressed (or not) will bear distinctive behavioral markers in their postings. A measure of overall engagement of author of a post in social media is *volume*, defined as the number of posts the user has made so far on social media. We define a second engagement feature to be the proportion of *reply* posts (@-replies) from a post’s author. The third feature is the fraction of *retweets* from a post author. The proportion of *links* (urls) shared by each post author comprises our fourth feature. We define a fifth feature as the fraction of *question-centric* posts from a post’s author.

Ego-network. We define two features that characterize a post author’s egocentric social network: (1) the number of *followers* or inlinks of the user, (2) the count of her *followees* or outlinks.

EXPLORATION OF DEPRESSION BEHAVIOR

Based on the features discussed so far, we present some descriptive analyses of differences in the two classes of posts: the “depression-indicative”, and the “standard” posts.

Table 3 gives a list of the high frequency unigrams that appear in the postings of the two classes. For the negative class (standard posts), most of the unigrams relate to commonplace details of daily life, ranging from work to entertainment (*work*, *friends*, *life*, *tomorrow*, *movie*, *football*). A number of positive emotion words are also observed, such as *brilliant*, *love*, *beautiful*, *perfect*, *great*.

Depression-indicative posts	Standard posts
loser, depress*, lonely, sad, alone, weak, useless, life, imbalance, blame, problems, unsuccessful, suicidal, torture, safe, escape, worry, intimidat*, uncomfortable, therapy, medication, shit, pressure, conversation, hurts, myself, worth, break, nobody, mine, painful, hate, suck*	lol, work, weekend, say, friends, brilliant, follow, tips, bieber, love, amazing, hello, now, bored, awesome, beautiful, romantic, fuck*, perfect, excited, smile, meet, tonight, life, movie, football, favorite, sleepy, great, night, team, good, anyone, you, your, tomorrow, money

Table 3. Top unigrams for the two classes of posts.

On the other hand, in the case of depression-indicative posts, many words are emotional in nature (e.g., *sad*, *happiness*, *uncomfortable*, *hurts*, *painful*, *hate*, *hope*, *worry*). However we notice a strong inclination towards unigrams representing negative affect and low intensity emotions—possibly reflecting the mental instability and helplessness of the individuals sharing these posts, including symptoms of their likely depression (*loser*, *depress**, *weak*, *useless*, *suicidal*, *unsuccessful*). Some posts also contain references to therapy and medication, possibly because individuals may be interested to exchange information about these with their audiences.

We extend these findings through observations of the difference in means across the two classes in terms of various features (except the *n*-grams)—for a particular feature *f*, it is defined as $\mu(f)=f_d - f_s$ where f_d (or f_s) is the

mean value of f for the positive (or negative) class. This is shown in Table 4.

There are considerable differences across the two classes, as noted in the differences of means as well as the statistical significance tests. For instance, for the depression-indicative posts, we observe considerable decrease in user engagement features, such as volume, RTs, and replies. This indicates that the authors of the depression-indicative posts are posting less, suggesting a possible loss of social connectedness. Additionally, lowered numbers of followers and followees shows that the authors of the depression-indicative posts exhibit reduced desire to socialize or tendency to consume external information and remain connected with others.

measure	<i>t</i> -stat	diff(μ)	measure	<i>t</i> -stat	diff(μ)
PA	14.64**	-0.031	func. words	3.72	-0.007
NA	16.03***	0.0295	article	9.75 *	0.156
activation	19.4 ***	-1.784	tentative	2.74	0.021
dominance	20.2 ***	-1.292	certainty	4.53 **	-0.059
time	15.81**	624.04	inhibition	7.3 *	0.035
1 st pp.	5.09 ***	0.155	inclusive	2.97	0.013
2 nd pp.	5.24 ***	-0.162	exclusive	4.83	-0.008
3 rd pp.	4.01 ***	-0.189	assent	5.16 *	-0.075
indef. pp.	13.24*	0.034	nonfluency	7.05 *	0.142
verbs	6.03 *	-0.048	filler	4.91 *	0.068
aux-verbs	4.48 *	-0.026	volume	15.21 ***	188.39
adverbs	6.13 *	0.087	RT	14.05**	90.25
preposition	3.91	-0.025	replies	22.88***	81.83
conjunction	1.23 **	0.117	links	14.49**	77.14
negation	8.42 *	0.092	questions	8.205*	73.95
quantifier	4.81	-0.037	#followers	10.6 **	-164.7
swear	12.91**	0.082	#followees	28.05***	-291.4

* $p < 0.01$; ** $p < .001$; *** $p < .0001$

Table 4. Difference of means comparing the two classes of posts over the different features discussed in previous section (except n -grams). Results of statistical significance tests, based on independent sample t -tests are also shown ($df=69,512$).

There is also high NA characterizing these post authors—possibly reflecting their mental instability and helplessness. Moreover, low activation and dominance may indicate loneliness, restlessness, exhaustion, lack of energy, and sleep deprivation, all of which are known to be consistent depression symptoms [6,18,23]. Finally, we find that the presence of the first-person pronoun is considerably high in posts of the same class, reflecting the users’ high attention to self and psychological distancing from others [3,6].

We weave together these observations, and the outcomes of the t -tests demonstrating statistically significant differences across the classes (see p -values in Table 4). It appears that our choice of features can adequately capture the distinctions across the depression-indicative and other posts. In this light, in the following section, we present a prediction framework that can classify a given post into one of the two classes, leveraging signal from these features.

PREDICTING DEPRESSION-INDICATIVE POSTS

Classification Framework

We use a supervised learning setup for classifying whether or not a given post is depression-indicative. We represent Twitter posts as vectors of the features presented earlier (e.g., emotion, time, n -grams, style, engagement features). Before feature extraction, the posts are lowercased, and numbers are normalized into a canonical form. Finally the posts are tokenized. After feature extraction, features that occur fewer than five times are removed in a first step of feature reduction. We then randomly split the data into five folds for cross-validation. The high dimensionality of the feature space can lead to overfitting to the training data—therefore we deploy principal component analysis (PCA) [7]. The classification algorithm is a standard Support Vector Machine classifier with an RBF kernel [7], although we experimented with other parametric and non-parametric supervised learning methods. We use five-fold cross validation, and conduct 100 randomized experimental runs.

Prediction Performance

Using the above model, we examine prediction performance in identifying the two classes of posts. In order to understand the importance of various feature types, we trained one model each using: (1) engagement and ego-network features; (2) n -grams; (3) linguistic style; (4) emotion and time features; (5) all features; and (6) dimensionality-reduced set of features.

	precision	recall	acc. (+ve)	acc. (mean)
eng. + ego.	0.624	0.617	56.159%	59.309%
n -grams	0.631	0.639	58.245%	60.002%
style	0.758	0.665	60.557%	65.758%
emo. + time	0.803	0.667	65.776%	68.648%
all features	0.826	0.674	69.740%	70.078%
dim. reduced	0.828	0.675	73.579%	74.576%

Table 5. Performance metrics in depression prediction in posts using various models. Third column shows the mean accuracy of predicting the positive class i.e., depression-indicative posts.

We present the results of these prediction models in Table 5. The results indicate that, in our test set, the best performing model (dimension-reduced features) yields an average accuracy of ~73% and high precision of 0.82, corresponding to the class of depression-indicative posts. Good performance of this classifier is also evident from the receiver-operator characteristic (ROC) curves in Figure 2. We find that the dimension reduced feature model gives slightly greater traction in prediction compared to the model that uses all features—demonstrating that reducing feature redundancy is important.

We also observe better performance for models that uses the linguistic style features alone, and the one that uses emotion and time features. Results in prior literature suggest that use of linguistic styles such as pronouns and articles provide information about how individuals respond to psychological triggers [3]. We conjecture that the general

social and psychological distancing characterizing the circumstances of clinically depressed individuals is linked to high attentional focus on oneself—as indicated by use of these styles (see also Table 4). Hence style features turn out to be a strong predictor of depression.

Next, the relatively better performance of the model using emotion and time features shows that expression of PA, NA as well as activation and dominance are central, so is the pattern of posting through the course of a day. Note that, as discussed earlier, one of the main characteristics of depression is disturbed processing of emotional information as indexed by disturbed startle reflex modulation, as well as a reduced sense of arousal in day-to-day activities [1,18]. Hence we observe better performance of these features in the prediction task. Similarly, psychiatric literature on depression indicates that 8 out of 10 people suffering from depression tend to worsen their symptoms during night [15]. Night time Internet activity is a known characteristic of these individuals [15]. This explains why the time feature, capturing the timestamp of post, is able to help predict if the post could be depression-indicative.




Figure 2. ROC curves showing performance of various prediction models.

In essence, we conclude that social media activity provides useful signals that can be utilized to classify and predict whether a post could be indicative of depression.

SOCIAL MEDIA DEPRESSION INDEX

Can the ability to predict whether or not a Twitter post is depression-indicative provide the basis for an accurate, reliable model of real-world depression rates in large populations? To this end, we use our predictive model, that considers whether a given post is depression-indicative, to *automatically label* a large corpus of posts shared on Twitter on any given day. Thereafter, we define a metric called the *social media depression index* (SMDI). This

index gives a measure of the degree of depression as manifested by Twitter users in their daily postings. On a given day t , we define it as the standardized difference between the frequencies of depression-indicative posts $n_d(t)$ and “standard” (or non-depression indicative) posts $n_s(t)$:

$$SMDI(t) = \frac{n_d(t) - \mu_d}{\sigma_d} - \frac{n_s(t) - \mu_s}{\sigma_s}, \quad (1)$$

where μ_d (correspondingly μ_s) and σ_d (correspondingly σ_s) are the mean and standard deviations of the number of depression-indicative (correspondingly standard) posts shared in a fixed time period before t (here between k and $t-1$, such that $1 \leq k \leq t-1$). Note that we consider separate terms for depression-indicative posts and standard posts. This allows depression and non-depression expressions in posts to be weighted equally (since their relative volumes are likely to be different—see, e.g., Table 1; also some days may be more “depression-inducing” than others). By standardizing depression and non-depression expression separately, we also focus on variation in each class separately. That is, even if per one’s behavior, some individuals dramatically under-express depression in their postings, each day’s relative depression expression compared to non-depression expression will be informative as a measure. In the above equation, note that $SMDI(t)$ is zero when there are as many standardized depression-indicative posts as non-depression indicative ones; while it will be positive for high depression days and negative for low depression ones. We note that Kramer [14] used a similar formulation while using Facebook for computing the Gross National Happiness index.

We also define an individual-centric metric of SMDI: Given a user u , we define her SMDI on day t as:

$$SMDI(u, t) = \frac{n_d(u, t) - \mu_d(t)}{\sigma_d(t)} - \frac{n_s(u, t) - \mu_s(t)}{\sigma_s(t)}, \quad (2)$$

where $n_d(u, t)$ (correspondingly $n_s(u, t)$) is the number of depression-indicative (correspondingly standard) posts from user u on day t , and $\mu_d(t)$ (correspondingly $\mu_s(t)$) and $\sigma_d(t)$ (correspondingly $\sigma_s(t)$) are the mean and standard deviations of the number of depression-indicative (correspondingly standard) posts shared on the same day t .

POPULATION CHARACTERISTICS OF DEPRESSION

We now utilize SMDI to understand a variety of population characteristics of depression.

Geographical Analysis

We present analyses on how our social media depression index reflects the known depression rates in several US cities and states. First we collected a list of 20 cities reported by businessweek.com to be the “unhappiest US cities”. The list includes a per-city depression rank, as measured by prescription drug claims for common

antidepressants². These statistics were reported at end of 2011. Using Twitter's Firehose, we crawled 30% random samples of Twitter posts (English language), per month between Jan 1 2011 and Dec 31 2011, for each city, to approximately match the time period around which these numbers were reported. Note that matching a post with a city was made based on the post author's self-report location on the Twitter profile, although other sophisticated location inference schemes may be adopted in future work. Table 6 gives statistics of the data crawled for each city.

City	Posts	Users	City	Posts	Users
Portland	3.38M	24K	Cleveland	6.95M	28K
Jacksonville	2.52M	17K	Milwaukee	3.48M	18K
Nashville	4.63M	23K	Sacramento	2.57M	17K
Seattle	5.87M	39K	Kansas City	3.15M	18K
Cincinnati	3.05M	17K	Tucson	1.69M	12K
Louisville	2.41M	16K	New Orleans	8.16M	23K
Minneapolis	2.69M	21K	Atlanta	26.1M	91K
Pittsburgh	5.41M	25K	Memphis	6.79M	27K
Indianapolis	4.49M	23K	Las Vegas	6.39M	37K
St Louis	421K	2.3K	Detroit	13.3M	43K

Table 6. Statistics of 20 US cities with high depression rate.

We used our trained prediction model to label all the posts shared every day from each city on whether they were depression-indicative, and thereafter computed the corresponding SMDI per day (eqn. (2)). We further calculated a mean SMDI per city throughout the time of analysis. Figure 3 shows the social media depression indices of all the 20 cities (y-axis) against their reported




Figure 3. Social media depression index of 20 US cities compared with their respective reported depression rates. A linear trend shows high correlation.

² Unhappy cities in the US:
http://images.businessweek.com/ss/09/02/0226_miserable_cities/index.htm

ranking in terms of depression rates (x-axis). The main observation in the figure is that we observe reasonable correlation ($R^2=0.64$) between our computed SMDI and known depression rates in the various cities. Note that, the slope of the trend line in the figure is positive (0.082); this is because we are predicting depression rank, for which lower rank implies worse condition.

On examination of the trend, it appears that we may be overestimating the degree of depression in more cities (e.g., Portland, Cleveland, Louisville, Atlanta, Kansas City), than the number of cities for which we underestimating the degree of depression (e.g., New Orleans, Detroit). To the extent this is true, this may be an artifact of a general tendency for people to express relatively more negativity on Twitter. Prior literature shows that negative affect is more predominantly expressed on Twitter than is positive affect [5]. Thus our model may be misinterpreting some expression of negativity as depression-indicative.

Next we expand our geographical analysis to understand how SMDI reflects actual depression rates in the 50 US states. In this case, we obtain data on depression rates in the various states from the Centers for Disease Control and Prevention (CDC), per the Behavioral Risk Surveillance

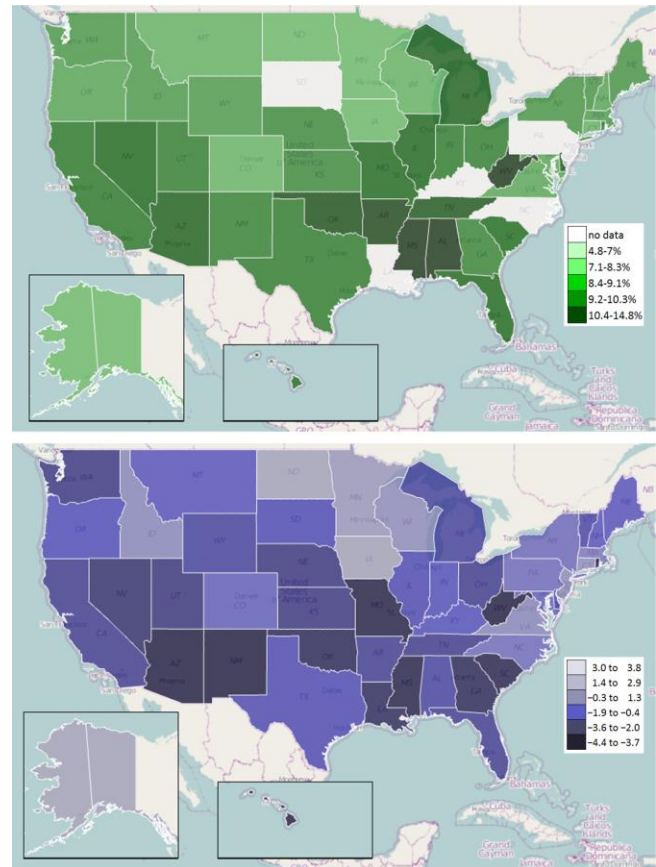


Figure 4. Heat map rendering of actual (CDC data) and predicted (SMDI) depression in various US states. Note that in both figures, higher intensity colors imply greater depression. A linear regression fit between the actual and predicted rates shows positive correlation of 0.5073.

Survey (BRFSS) conducted in 2008 [2]. For computing SMDI, we follow a similar approach as done with the cities. That is, we compile a set of Twitter posts for each state based on regex matching of the state's name with the user's self-reported location on Twitter. Our Twitter data go back to June 2010, and thus to best align with the 2008 CDC data, we collected a 30% sample of Twitter data over a six month period between Jun 1, 2010 and Dec 31, 2010. There were a total of approximately 132M posts over all the 40 states in our sample thus compiled.

Figure 4 shows heat map rendering of actual and predicted depression rates in various states—higher intensity colors imply greater depression rate reported by CDC (top) or greater SMDI (bottom). Using the exact percent of each state's population that is depressed (as opposed to the binned values in Figure 3), the Pearson correlation between actual per-state rate of depression and those predicted by SMDI is found to be 0.51. This demonstrates that our metric can capture to a reasonable extent, the actual rates of depression known in these states.

Demographic Analysis: Gender Differences

In this section we perform a demographic-centric analysis of depression, comparing men and women. Our goal is to be able to examine whether the rates of depression as found by the BRFSS survey (same source as above) for men and women align with the SMDI rates.

For the purpose, we utilize the same corpus of Twitter posts used in the analysis reported in Figure 4. This analysis however necessitated that we know the gender of the author of a post, which is not an available attribute from Twitter. Hence we utilized a gender classifier [6] that uses regex matching of Twitter's self-reported first names of users with a gazetteer of names and gender collated from US Census, baby name lists and public Facebook profiles. As found in [6], a gender classifier of this type is known to yield accuracy between 85-90%. In our corpus, we obtained




Figure 5. Differences in actual and predicted depression for men and women based on Twitter posts, and from users based in the US.




Figure 6. Diurnal patterns of social media depression index for men and women.

about 61% women and the rest men, after about 65% coverage over all users.

For each post for which the author's gender was inferred, we computed the individual-centric measure of SMDI (see eqn. (3)). Figure 5 shows the results. We observe that we can reasonably predict the percentages of men and women depressed. Women are known to suffer from 1.3 times more depression than men [11]—approximately that ratio is observed in our prediction as well (1.5 times in our case as given by the actual rates of percent depressed in Figure 5). This is denoted by the nearly parallel lines for predicted and actual percent of each gender who are depressed. However as with our previous analysis on geography, we observe a slight overestimation of depression when using our method. Interestingly, it appears that overestimation is higher for women than for men (notice the SMDI values specifically). One possibility is that women generally show greater emotional expressivity than men, typically negative in affect [8], hence leading to a bias in our model.

Temporal Analysis

In this final section on population-scale analyses, we focus on understanding some of the temporal patterns of depression, at diurnal and annual scales.




Figure 7. Yearly patterns of social media depression index over US cities with high depression rates.

We first use the same data as used in Figure 5 to study levels of SMDI over the course of a typical day (in local time) across men and women. This is shown in Figure 6. As in correspondence with our findings in Figure 5, we see that diurnal SMDI values for women ($\mu=-0.47$) are higher than that of men ($\mu=-0.139$; $p<.001$); although they have high correlation ($R^2=0.78$). An interesting finding of this study is that, for both men and women, SMDI is higher during the night than during the day, with SMDI values peaking at around midnight, and the lowest value appearing around noon. As mentioned before, psychiatric literature on depression indicates that 8 out of 10 people suffering from depression exhibit worsened symptoms at night [15]: People become more prone to worrying thoughts at night due to loneliness, break from work, lack of energy, or other interactions between light/darkness and the nervous system.

Next, we study the rhythms of SMDI over the course of a year, aggregated in pit-stops of months. For this purpose, we expand the above dataset, and derive two 30% samples of Twitter posts ranging between Jun 2010-May 2011, and Jan 2012-Dec 2012 (307M posts in all). The yearly patterns of SMDI for several US cities (chosen for exemplary purposes) with high depression rate (see Figure for details) are shown in Figure 7 (averaged across the two dataset snapshots). We also show an average trend for all posts, combining all locations (denoted as “All”).

We observe that the yearly depression pattern shows a seasonal trend, with highest depression observed during winter season in the US ($\mu=0.47$; Dec-Jan), while the lowest during summer ($\mu=-1.88$; Jun-Aug) and to some extent fall ($\mu=-0.73$; Sep-Nov). Clinical literature indicates that there is greater prevalence of depressive anxiety during winter at northern latitudes, sometimes attributed to insufficient exposure to light during winter [25].

Across the cities, we observe that SMDI in certain cities is more associated with weather conditions than others. Jacksonville and Seattle are both ranked high in terms of depression rates, however the variation in SMDI trend for Seattle ($\sigma^2=4.45$) is much higher than that for Jacksonville ($\sigma^2=0.85$). In fact, the percent difference between Seattle and Jacksonville’s SMDI during winter is 8% higher than that during summer. Note that Seattle’s seasonal weather variations are more extreme than those for Jacksonville, per National Oceanic and Atmospheric Administration (NOAA). As also supported by clinical literature, we thus conjecture that Twitter users based in Seattle are more prone to depressive symptoms during winter than in Jacksonville, or other low weather variability cities.

However, notice that even during the summer, cities like Jacksonville and other low weather variation cities (e.g., Nashville) show considerable depression prevalence—SMDI for Jacksonville in summer is $\mu=-1.3$, which is higher compared to some of the other cities. That is, it appears that while weather may have a strong role to play in the depression rates, there are likely to be other factors at play,

for instance, socio-economic status, unemployment rates, which may explain the higher SMDI in these cities.

To summarize, we observe persistent rhythms in depression expression on social media during the course of a day across men and women, and even at a seasonal level across a variety of locations. This provides us with a promising mechanism to monitor fine-grained temporal trends of depression across populations, demographics, and regions.

DISCUSSION

Implications

Through our experimental findings, we have demonstrated how sets of behavioral markers manifested in social media can be harnessed to predict depression-indicative postings, and thereby understand large-scale depression tendencies in populations. Since our predictions can be made considerably more frequently than BRFSS or CDC surveys, the resultant population-scale estimates of depression can be utilized time to time to enable early detection and rapid treatment of depression. Moreover, a primary challenge in public health is that several behavioral health concerns go under-reported—we hope that through our proposed technique, organizations may be able to gauge better such concerns and improve healthcare support.

Note that our method of depression prediction, or the analyses are, however, not meant to be a replacement for traditional surveillance systems, or clinical diagnoses of depression. Rather, the population-scale trends over geography, time or gender may be a mechanism to trigger public health inquiry programs to take appropriate and needful measures, or allocate resources in a timely manner.

Limitations

Although the analysis of social media postings makes it possible to track depression levels in ways there are not feasible offline, there is an inherent population bias in our study. According to Pew Research Center [21], among the 74% of American adults who use the internet, only about 8% report using Twitter. Additionally, unlike surveys, we have little knowledge about people’s idiosyncratic behavior “behind the scenes”, their social, cultural and psychological environment, or socio-economic status. Potentially, the limitations of Twitter may be tackled in one way by adding complementary sources of behavioral data, such as about social ties from Facebook, web browsing behavior, or search query logs, in conjunction with health records, such as antidepressant purchases, or healthcare claims data. These opportunities remain ripe areas of future research.

Privacy

This research revolves around analysis individuals’ behavioral health and leverages information that may be considered sensitive, hence privacy is important. In all of our datasets, as well as the initial crowdsourcing study, we did not retain any information about users’ real identity (if available). In the Mechanical Turk task, individuals could opt out, if they wished to, from sharing their Twitter

information; and those who did were informed that their data may be used for research purposes. It is also important to also note that we make use of *public* Twitter postings exclusively, and our predictions are made at the level of posts, instead of individuals, preventing any conspicuous associations that may arise between the observed social media activity of an individual and their actual psychological state.

CONCLUSION

Mining and analysis of social media activity in order to understand a variety of public health phenomena has been gaining considerable traction recently among researchers. In this paper, we have demonstrated the potential of using social media as a reliable tool for measuring population-scale depression patterns. We adopted a crowdsourcing strategy of collecting ground truth data on depression from Twitter, and devised a variety of measures such as language, emotion, style and user engagement to build an SVM classifier. The classifier predicted with high accuracy (73%) whether or not a post on Twitter could be depression-indicative. Thereafter, the trained model was leveraged in a population-scale measurement metric of depression—called the social media depression index. Variety of analyses around geography, gender and time showed that SMDI can closely mirror CDC defined statistics on depression. In the future, we are interested in developing individual-centric predictive models that analyze a person's social media feeds, and provide early warning/intervention if there are behavioral concerns out of the ordinary. Modeling the contagion of depressive disorders in social media is also an exciting future direction.

REFERENCES

- Billings, A., Moos, Rudolf H. (1984). Coping, stress, and social resources among adults with unipolar depression. *Journal of Personality and Social Psychology*, 46(4), 877-891.
- Centers for Disease Control and Prevention (CDC). *Behavioral Risk Factor Surveillance System Survey Data*. Atlanta, Georgia: U.S. Department of Health and Human Services, Centers for Disease Control and Prevention, 2008, 2011, 2012.
- Chung, C.K., & Pennebaker, J.W. (2007). The psychological functions of function words. In K. Fielder (Ed.), *Social communication* (pp. 343-359).
- Collier, N., Son, N., & Nguyen, N. (2011). OMG U got flu? Analysis of shared health messages for bio-surveillance. *Journal of Biomedical Semantics*.
- De Choudhury, M., Counts, S., and Gamon, M. (2012). Not All Moods are Created Equal! Exploring Human Emotional States in Social Media. In *Proc. ICWSM '12*.
- De Choudhury, M., Counts, S., & Horvitz, E. (2013). Predicting Postpartum Changes in Behavior and Mood via Social Media. In *Proc. CHI 2013*, to appear.
- Duda, Richard O., Hart, Peter E., & Stork, David G. (2000). *Pattern Classification*. 2nd Edition, Wiley.
- Fujita, F., Diener, E. and Sandvik, E. (1991). Gender differences in negative affect and wellbeing: The case for emotional intensity. *Journal of Personality and Social Psychology*, 61: 427-434.
- González HM, Vega WA, Williams DR, Tarraf W, West BT, & Neighbors HW (2010). Depression Care in the United States: Too Little for Too Few. *Archives of General Psychiatry* 67 (1): 37-46.
- Kawachi, I., and Berkman, L. S. (2001). Social ties and mental health. *Journal of Urban Health*, 78(3), 458-467.
- Kessler, R.C., Berglund, P., Demler, O. et al. (2003). The Epidemiology of Major Depressive Disorder: Results from the National Comorbidity Survey Replication (NCS-R). *Journal of the American Medical Association* 289 (23): 3095-3105.
- Kramer, A. (2010). An Unobtrusive Behavioral Model of "Gross National Happiness". In *Proc. CHI 2010*.
- Krieck, M., Dreesman, J., Otrusina, L., & Denecke, K. (2011). A new age of public health: Identifying disease outbreaks by analyzing tweets. In *Proc. Health Web-Science Workshop*, ACM Web Science Conference.
- Luoma JB, Martin CE, Pearson JL (2002). Contact with mental health and primary care providers before suicide: a review of the evidence. *The American Journal of Psychiatry* 159 (6): 909-16.
- Lustberg L, & Reynolds CF (2000). Depression and insomnia: questions of cause and effect. *Sleep Medicine Reviews* 4 (3): 253-262.
- Mathers, C.D., Loncar, D. (2006). Projections of global mortality and burden of disease from 2002 to 2030. *PLoS Med* 3 (11).
- Niels Rosenquist, J., Fowler, J. & Christakis, N. (2011). Social Network Determinants of Depression. *Molecular Psychiatry* 16 (3): 273-281.
- Oxman T.E., Rosenberg S.D., & Tucker G.J. (1982). The language of paranoia. *American J. Psychiatry* 139:275-82.
- Park, M., Cha, C., & Cha, M. (2012). Depressive Moods of Users Captured in Twitter. In *Proc. ACM SIGKDD Workshop on Healthcare Informatics (HI-KDD)*.
- Paul, M., J., & Dredze, M. (2011). You are What You Tweet: Analyzing Twitter for Public Health. In *Proc. ICWSM '11*.
- Pew Research Center Report on Twitter use: <http://pewinternet.org/Reports/2010/Twitter-Update-2010/Findings/Overview.aspx>
- Radloff, L.S. (1977). The CES-D scale: A self-report depression scale for research in the general population. *Applied Psychological Measurement* 1: 385-401.
- Rude, S., Gortner, E., & Pennebaker, J. (2004). Language use of depressed and depression-vulnerable college students. *Cognition and Emotion*, 1121-1133.
- Sadilek, A., Kautz, H., & Silenzio, V. (2012). Modeling Spread of Disease from Social Interactions. In *Proc. ICWSM '11*.
- Wehr TA, Rosenthal NE. (1989). Seasonality and Affective Illness. *Am J Psychiatry*. 146(7):829-39.